

Problem Set #1 Sample Solutions

Greg Plaxton

February 11, 2019

1(a). Write ℓ as $2^k + r$, where $0 \leq r < 2^k$. We claim that an ℓ -leaf binary tree with minimum external path length has $2r$ leaves at depth $k + 1$ and $\ell - 2r$ leaves at depth k (the root is at depth 0). We prove this claim in two stages. First, we argue that any ℓ -leaf binary tree with minimum external path length is *full*, that is, every internal node has degree two. This is easy to see since the external path length of a non-full ℓ -leaf binary tree can be strictly reduced by removing a degree-one internal node, replacing it with its child. We now complete the proof of the claim by arguing that any full ℓ -leaf binary tree with minimum external path length has all of its leaves on at most two adjacent levels. To see this, consider a full ℓ -leaf binary tree with deepest leaf x at depth d and shallowest leaf y at level d' where $d' < d - 1$. Because the tree is full and x is a deepest leaf, x has a sibling leaf x' , also at depth d . Let p denote the parent of leaves x and x' . Consider modifying this tree by making x and x' children of y instead of p . The resulting tree still has ℓ leaves: The only changes to the set of leaves are that p has been added and y has been removed. The change in the external path length is easy to calculate: There is a decrease of $2(d - d' - 1)$ due to the change in the depths of x and x' , and there is an increase of $d - d' - 1$ due to the replacement of y by p in the set of leaves, for a net decrease of $d - d' - 1$, which is a positive quantity. This completes the proof of the claim. The desired $\Omega(\ell \lg \ell)$ lower bound now follows easily.

An alternative proof uses (strong, i.e., course of values) induction on ℓ , where the key technical observation needed to carry out the induction step is that the function $f(x) = x \lg x$ is convex (positive second derivative), so that an expression such as $f(a) + f(\ell - a)$ is minimized when $a = \frac{\ell}{2}$.

1(b). Fix a comparison-based sorting algorithm A . Algorithm A corresponds to a decision tree T where each of the $n!$ possible input permutations descends to a distinct leaf of T . (If two or more permutations descend to the same leaf, then A is not a correct sorting algorithm.) For a given permutation π , the depth of the leaf of T to which π descends corresponds to the number of comparisons made by A on input π . Thus the external path length of T is at least the sum, over all $n!$ permutations π , of the number of comparisons made by A on input π . In other words, the external path length of T is at least $n!$ times the average number of comparisons used by A . The result now follows from part (a) and the fact that $\lg(n!) = \Theta(n \lg n)$.

2(a). It is easy to check that $n_0 = 5$. Let n_1 denote the least positive integer such that $n \geq 4\sqrt{n} \ln n + 8$ holds for all $n \geq n_1$. It is easy to check that $n_1 \geq n_0$. Let c denote

$$\max \left(2, \max_{2 \leq n < n_1} \frac{T(n)}{n \ln n} \right)$$

Let $P(n)$ denote the predicate “ $T(n) \leq cn \ln n$ ”. We claim that $P(n)$ holds for all $n \geq 2$. The definition of c implies that $P(n)$ holds for $2 \leq n < n_1$. Fix an integer $n \geq n_1$ and assume inductively that $P(k)$ holds for $2 \leq k < n$. To see that $P(n)$ holds, observe that

$$\begin{aligned} T(n) &= 2T\left(\left\lfloor \frac{n}{2} + \sqrt{n} \right\rfloor\right) + n \\ &\leq 2c \left\lfloor \frac{n}{2} + \sqrt{n} \right\rfloor \ln \left\lfloor \frac{n}{2} + \sqrt{n} \right\rfloor + n \\ &\leq 2c \left(\frac{n}{2} + \sqrt{n} \right) \ln \left(\frac{n}{2} + \sqrt{n} \right) + n \\ &= 2c \left(\frac{n}{2} + \sqrt{n} \right) \left[\ln \frac{n}{2} + \ln \left(1 + \frac{2}{\sqrt{n}} \right) \right] + n \\ &\leq 2c \left(\frac{n}{2} + \sqrt{n} \right) \left(\ln n - 1 + \frac{2}{\sqrt{n}} \right) + n \\ &= cn \ln n - cn + 2c\sqrt{n} + 2c\sqrt{n} \ln n - 2c\sqrt{n} + 4c + n \\ &= cn \ln n - (c-1)n + 2c\sqrt{n} \ln n + 4c \\ &\leq cn \ln n - (c/2)n + 2c\sqrt{n} \ln n + 4c \\ &= cn \ln n - (c/2)(n - 4\sqrt{n} \ln n - 8) \\ &\leq cn \ln n, \end{aligned}$$

where the first equality follows from the definition of $T(n)$, the first inequality follows from the induction hypothesis (note that $2 \leq \frac{n}{2} + \sqrt{n} < n$ since $n \geq n_1 \geq n_0$), the third inequality follows from $\ln(1+x) \leq x$ for $x > -1$, the fourth inequality follows from $c \geq 2$, and the last inequality holds since $n \geq n_1$ implies $n \geq 4\sqrt{n} \ln n + 8$.

2(b). $f(n) = \frac{n}{\ln n}$.

3(a). Let $B(x)$ and $C(x)$ denote the quotient and remainder polynomials when $A(x)$ is divided by $x - z$. Since $x - z$ has degree 1, $C(x)$ has degree zero, i.e., $C(x)$ is of the form $c_0 x^0 = c_0$. We have $A(x) = B(x) \cdot (x - z) + C(x) = B(x)(x - z) + c_0$, and thus $A(z) = c_0$, as required.

3(b). We have $P_{k,k}(x) = (x - x_k)$ and hence $Q_{k,k}(x) = A(x) \bmod (x - x_k)$. Thus the result of part (a) implies $Q_{k,k}(x) = A(x_k)$.

It remains to prove that $Q_{0,n-1}(x) = A(x)$. We have $P_{0,n-1}(x) = \prod_{0 \leq k < n} (x - x_k)$. Since $A(x)$ has degree at most $n - 1$ and $P_{0,n-1}(x)$ has degree n , we deduce that $A(x)$ is the remainder when $A(x)$ is divided by $P_{0,n-1}(x)$. Thus $Q_{0,n-1}(x) = A(x)$.

3(c). Fix integers i, j , and k such that $0 \leq i \leq k \leq j < n$. Let $B(x)$ denote the quotient polynomial when $A(x)$ is divided by $P_{i,j}(x)$. Thus $A(x) = B(x) \cdot P_{i,j}(x) + Q_{i,j}(x)$ where $P_{i,j}(x)$ has degree $j - i + 1$ and $Q_{i,j}(x)$ has degree at most $j - i$. Let $C(x)$ and $D(x)$ denote the quotient and remainder polynomials when $Q_{i,j}(x)$ is divided by $P_{i,k}(x)$. Thus $Q_{i,j}(x) = C(x) \cdot P_{i,k}(x) + D(x)$ where $P_{i,k}(x)$ has degree $k - i + 1$ and $D(x)$ has degree at most $k - i$. It follows that

$$\begin{aligned} A(x) &= B(x) \cdot P_{i,j}(x) + C(x) \cdot P_{i,k}(x) + D(x) \\ &= (B(x) \cdot P_{k+1,j}(x) + C(x)) P_{i,k}(x) + D(x). \end{aligned}$$

Let $E(x)$ denote the polynomial $B(x) \cdot P_{k+1,j}(x) + C(x)$. Since $A(x) = E(x) \cdot P_{i,k}(x) + D(x)$ where $P_{i,k}(x)$ has degree $k - i + 1$ and $D(x)$ has degree at most $k - i$, we conclude that $E(x)$ and $D(x)$ are the quotient and remainder polynomials when $A(x)$ is divided by $P_{i,k}(x)$. Thus $D(x) = Q_{i,k}(x)$, as required.

A symmetric argument shows that $Q_{k,j}(x)$ is the quotient polynomial when $Q_{i,j}(x)$ is divided by $P_{k,j}(x)$.

3(d). Assume for simplicity that n is a power of 2. We find it easiest to describe our algorithm in terms of a complete binary tree T with n leaves indexed from 0 to $n - 1$. (Such a tree can be implemented as an array.) For any node u of the tree, let $u.a$ and $u.b$ denote the minimum and maximum indices of the leaves in the subtree of T rooted at u . We perform two passes over the tree. The first pass is upward — from the leaves to the root — and the second pass is downward — from the root to the leaves.

The goal of the first pass is to compute the coefficients of the polynomial $P_{u.a,u.b}(x)$ at each node u in T . For a leaf node u , this can be done in $O(1)$ time. Now consider an internal node u with children u_L and u_R . When we come to process node u , nodes u_L and u_R have already been processed, so the coefficients of the two polynomials $P_{u_L.a,u_L.b}(x)$ and $P_{u_R.a,u_R.b}(x)$ have already been determined. Since $P_{u.a,u.b}(x)$ is the product of these two polynomials, we can use the FFT algorithm to compute the coefficients of $P_{u.a,u.b}(x)$ in $O(s \log s)$ time where $s = u.b - u.a + 1$. Let $A(n)$ denote the total running time of the first pass. Thus $A(1) = O(1)$ and $A(n) = 2A(n/2) + O(n \log n)$ for $n > 1$. Thus $A(n) = O(n \log^2 n)$.

The goal of the second pass is to compute the coefficients of the polynomial $Q_{u.a,u.b}(x)$ at each node u in T . In the case where u is the root node, this is straightforward since $Q_{u.a,u.b}(x) = Q_{0,n-1}(x) = A(x)$ by the second claim of part (b). Now assume that node u is a proper descendant of the root, and let u_P denote the parent of u . When we come to process node u , node u_P has already been processed, so the coefficients of the polynomial $Q_{u_P.a,u_P.b}(x)$ have been determined. Let us assume that u is the left child of u_P ; a symmetric argument holds for the case where u is the right child of u_P . As a result of the first pass, the coefficients of the polynomial $P_{u.a,u.b}(x)$ are available at nodes u_P and u . Using the result of part (c), we have $Q_{u.a,u.b}(x) = Q_{u_P.a,u_P.b}(x) \bmod P_{u.a,u.b}(x)$. Thus, using the assumption stated in the problem concerning the complexity of polynomial division, we can compute the coefficients of the polynomial $Q_{u.a,u.b}(x)$ in $O(s \log s)$ time where $s = u_P.b - u_P.a + 1$. Let $B(n)$ denote the

total running time of the second pass. Thus $B(1) = O(1)$ and $B(n) = 2B(n/2) + O(n \log n)$ for $n > 1$. Thus $B(n) = O(n \log^2 n)$.

The first claim of part (b) implies that for each k in $\{0, \dots, n-1\}$, the value $A(x_k)$ is located at the leaf node with index k at the end of the second pass. Since $A(n)$ and $B(n)$ are each $O(n \log^2 n)$, the overall running time is $O(n \log^2 n)$, as required.

4(a). It is convenient to reindex the tasks from 1 to n in nondecreasing order of deadline, breaking ties arbitrarily.

For any integer j such that $0 \leq j \leq n$, let U_j denote the set of tasks $\{1, \dots, j\}$. For any set of tasks U , we define $\text{greedy}(U)$ as the schedule that executes the tasks in U in order of increasing index. For any set of tasks U , we define the *duration* of U as $\sum_{j \in U} e_j$, and we define the *value* of U as the value of $\text{greedy}(U)$. We say that a schedule S for a set of tasks U is feasible if every task in U meets its deadline in S . We say that a set of tasks U is feasible if there exists a feasible schedule for U .

We claim that a set of tasks U is feasible if and only if $\text{greedy}(U)$ is feasible. The “if” direction is immediate. To prove the “only if” direction, assume that U is feasible, and let S be a feasible schedule for U . If S is equal to $\text{greedy}(U)$, there is nothing further to prove. Otherwise, there exist distinct tasks i and j in U such that $i > j$ and task i is scheduled immediately before task j in S . It is easy to check that if we modify schedule S by interchanging the order in which tasks i and j are scheduled to obtain a new schedule S' , then S' is also a feasible schedule for U . Furthermore, the number of inversions in the permutation associated with S' is exactly one less than the number of inversions in the permutation associated with S . By repeating this argument, we eventually arrive at a feasible schedule for U with zero inversions. This feasible schedule is $\text{greedy}(U)$, completing the proof of the claim.

Let T denote $\sum_{1 \leq j \leq n} e_j$. For any integers i and j such that $0 \leq i \leq T$ and $0 \leq j \leq n$, we define $a_{i,j}$ as the maximum, over all feasible subsets U of U_j with duration exactly i , of the value of $\text{greedy}(U)$. (If no such subset U exists, we define $a_{i,j}$ as $-\infty$.) We now develop a recurrence for computing the $a_{i,j}$'s. In the following case analysis, let U be a feasible subset of U_j with duration exactly i and such that the value of $\text{greedy}(U)$ is $a_{i,j}$.

Case 1: $i = 0$. Then U is the empty set, and hence $a_{i,j} = 0$.

Case 2: $i > 0$ and $j = 0$. Then no such U exists, and hence $a_{i,j} = -\infty$.

Case 3: $i > 0$ and $j > 0$.

Case 3.1: $d_j < i$ or $i < e_j$. Then task j cannot belong to U , and hence $a_{i,j} = a_{i,j-1}$.

Case 3.2: $e_j \leq i \leq d_j$. If task j belongs to U , then $a_{i,j} = v_j + a_{i-e_j,j-1}$. If task j does not belong to U , then $a_{i,j} = a_{i,j-1}$. Thus $a_{i,j} = \max(a_{i,j-1}, v_j + a_{i-e_j,j-1})$.

The total number of operations required to compute the $a_{i,j}$'s using the above recurrence is $O(nT)$. The part (a) assumption implies that $T = O(n^{c+1})$. Thus the total number of operations required to compute the desired value $a_{i,j}$'s is $O(n^{c+2})$. The desired answer is given by $\max_{0 \leq i \leq T} a_{i,n}$.

4(b). Let V denote $\sum_{1 \leq j \leq n} v_j$. For any integers i and j such that $0 \leq i \leq V$ and $0 \leq j \leq n$, we define $b_{i,j}$ as the minimum, over all feasible subsets U of U_j with value exactly

i , of the duration of U . (If no such subset U exists, we define $b_{i,j}$ as ∞ .) We now develop a recurrence for computing the $b_{i,j}$'s. In the following case analysis, let U be a feasible subset of U_j with value exactly i and such that the duration of $\text{greedy}(U)$ is $b_{i,j}$.

Case 1: $i = 0$. Then U is the empty set, and hence $b_{i,j} = 0$.

Case 2: $i > 0$ and $j = 0$. Then no such U exists, and hence $b_{i,j} = \infty$.

Case 3: $i > 0$ and $j > 0$.

Case 3.1: $i < v_j$ or $d_j < e_j + b_{i-v_j,j-1}$. Suppose task j belongs to U . Then $i \geq v_j$, and hence $d_j < e_j + b_{i-v_j,j-1}$. The definition of U implies that $b_{i-v_j,j-1}$ is equal to the duration of the feasible subset $U \setminus \{j\}$ of U_{j-1} . Thus the inequality $d_j < e_j + b_{i-v_j,j-1}$ implies that task j misses its deadline in $\text{greedy}(U)$, contradicting the feasibility of U . We conclude that task j does not belong to U , and hence that $b_{i,j} = b_{i,j-1}$.

Case 3.2: $i \geq v_j$ and $e_j + b_{i-v_j,j-1} \leq d_j$. If task j does not belong to U , then $b_{i,j} = b_{i,j-1}$. If task j belongs to U , then $b_{i,j} = e_j + b_{i-v_j,j-1}$. Thus $b_{i,j} = \min(b_{i,j-1}, e_j + b_{i-v_j,j-1})$.

The total number of operations required to compute the $b_{i,j}$'s using the above recurrence is $O(nV)$. The part (b) assumption implies that $V = O(n^{c+1})$. Thus the total number of operations required to compute the desired value $b_{i,j}$'s is $O(n^{c+2})$. The desired answer is given by the maximum value of i for which $b_{i,n} < \infty$.