

# A Load Balancing Problem

- ▶ We are given  $n$  items to be placed in  $m$  bins
  - ▶ Item  $i$  has positive integer weight  $w_i$ ,  $1 \leq i \leq n$
  - ▶ We wish to assign each item to a bin in a way that minimizes the maximum bin load
  - ▶ Remark: The “load” of a bin is the total weight of the items assigned to it
- ▶ We will prove that the decision version of this problem, call it BALANCE, is NP-complete
  - ▶ Given a threshold  $B$ , is there a way to assign the items to the bins so that every bin has load at most  $B$ ?
  - ▶ It is easy to argue that BALANCE belongs to NP

# The Subset Sum Problem

- ▶ In the SUBSET SUM problem, we are given  $n$  integers  $x_1, \dots, x_n$  and an integer target  $T$ , and we wish to determine whether there is a subset  $I$  of  $\{1, \dots, n\}$  such that 
$$\sum_{i \in I} x_i = T$$
- ▶ It is easy to argue that SUBSET SUM belongs to NP
- ▶ The text uses a reduction from 3-SAT to prove that SUBSET SUM is NP-complete

# The Partition Problem

- ▶ In the PARTITION problem, we are given  $n$  positive integers  $x_1, \dots, x_n$ , and we wish to determine whether there is a subset  $I$  of  $\{1, \dots, n\}$  such that

$$\sum_{i \in I} x_i = \sum_{i \in \{1, \dots, n\} \setminus I} x_i = \frac{1}{2} \sum_{1 \leq i \leq n} x_i$$

- ▶ It is easy to see that PARTITION belongs to NP
- ▶ We will use a reduction from SUBSET SUM to prove that PARTITION is NP-complete

# A Reduction from SUBSET SUM to PARTITION

- ▶ Let  $X$  be an instance of SUBSET SUM with positive integers  $x_1, \dots, x_n$  and target  $T$ 
  - ▶ Let  $S$  denote  $\sum_{1 \leq i \leq n} x_i$
  - ▶ Since target  $T$  is effectively the same as target  $S - T$ , we can assume without loss of generality that  $T \geq S/2$
  - ▶ If  $T = S/2$  then  $X$  corresponds directly to an instance of PARTITION (simply drop the target  $T$ ), so we can assume that  $T > S/2$
- ▶ Let  $Y$  denote the PARTITION instance with  $n + 1$  positive integers  $x_1, \dots, x_{n+1}$  where  $x_{n+1} = 2T - S$
- ▶ We claim that  $X$  is a positive instance of SUBSET SUM if and only if  $Y$  is a positive instance of PARTITION

# “Only If” Direction

- ▶ Assume that  $X$  is a positive instance of SUBSET SUM
- ▶ Hence there is a subset  $A$  of  $\{1, \dots, n\}$  such that  $\sum_{i \in A} x_i = T$
- ▶ Let  $B$  denote  $\{1, \dots, n+1\} \setminus A$
- ▶ Thus  $\sum_{i \in B} x_i = (S - T) + (2T - S) = T$
- ▶ Thus  $Y$  is a positive instance of PARTITION

# “If” Direction

- ▶ Assume that  $Y$  is a positive instance of PARTITION
- ▶ Thus we can partition  $\{1, \dots, n+1\}$  into two sets  $A$  and  $B$  such that  $\sum_{i \in A} x_i = \sum_{i \in B} x_i = [S + (2T - S)]/2 = T$
- ▶ Assume without loss of generality that  $n+1$  belongs to  $B$
- ▶ Thus  $A$  is a subset of  $\{1, \dots, n\}$  such that  $\sum_{i \in A} x_i = T$
- ▶ Thus  $X$  is a positive instance of SUBSET SUM

# A Reduction from PARTITION to BALANCE

- ▶ Let  $X$  be an instance of PARTITION with positive integers  $x_1, \dots, x_n$ 
  - ▶ Let  $S$  denote  $\sum_{1 \leq i \leq n} x_i$
  - ▶ We can assume without loss of generality that  $S$  is even
- ▶ Let  $Y$  be the instance of BALANCE with item weights  $x_1, \dots, x_n$ ,  $m = 2$  bins, and threshold  $B$  equal to  $S/2$
- ▶ Observe that  $X$  is a positive instance of PARTITION if and only if  $Y$  is a positive instance of BALANCE
- ▶ Thus BALANCE is NP-complete, even in the special case  $m = 2$

# The Class of NP-Hard Problems

- ▶ We say that a problem  $X$  is NP-hard if existence of a polynomial-time algorithm for  $X$  implies that  $P = NP$ 
  - ▶ Importantly, the problem  $X$  need not be a decision problem
  - ▶ For example, the problem of determining a maximum independent set of a graph is NP-hard
- ▶ It is easy to see that the load balancing problem is NP-hard
- ▶ It is natural to seek approximation algorithms for such NP-hard optimization problems



# Approximation Algorithms for Optimization Problems

- ▶ We typically seek approximation algorithms achieving a good approximation ratio
  - ▶ For a minimization problem, we say that an algorithm achieves an approximation ratio of  $\alpha$  if it is guaranteed to produce a solution with objective function value at most  $\alpha$  times optimal
  - ▶ For a maximization problem, we say that an algorithm achieves an approximation ratio of  $\alpha$  if it is guaranteed to produce a solution with objective function value at least  $1/\alpha$  times optimal
  - ▶ Thus the approximation ratio is at least 1, and an optimal algorithm achieves an approximation ratio of 1

# Approximability of NP-Hard Optimization Problems

- ▶ Given an NP-hard optimization problem, we seek to design an approximation algorithm with an approximation ratio as close to 1 as possible
- ▶ As we have seen, many NP-complete decision problems have an associated NP-hard optimization version
  - ▶ It turns out that these optimization problems vary widely in terms of polynomial-time approximability
  - ▶ For example, in some cases we can give a polynomial-time algorithm that achieves an approximation ratio of  $1 + \varepsilon$  for any constant  $\varepsilon > 0$
  - ▶ In other cases we can show that for any  $\varepsilon > 0$ , no polynomial-time algorithm achieves an approximation ratio of  $n^{1-\varepsilon}$  unless  $P = NP$

# Greedy Approximation Algorithms

- ▶ For many NP-hard optimization problems, we can use a simple greedy algorithm to achieve a good approximation ratio
- ▶ Today we will analyze two greedy algorithms for the load balancing problem
- ▶ The first of these two algorithms, which we denote  $\mathcal{A}$ , processes each successive item (from 1 to  $n$ ) by placing it in a least-loaded bin
  - ▶ We will prove that algorithm  $\mathcal{A}$  achieves an approximation ratio of 2
- ▶ The second algorithm, denoted  $\mathcal{B}$ , is the same as  $\mathcal{A}$  except that it processes the items in nonincreasing order of weight
  - ▶ We will prove that algorithm  $\mathcal{B}$  achieves an approximation ratio of  $4/3$

# Algorithm $\mathcal{A}$ : An Upper Bound of 2

- ▶ Fix an instance of the load balancing problem, and let  $L^*$  denote the minimax bin load
- ▶ Suppose that at some step of algorithm  $\mathcal{A}$ , we add an item  $i$  with weight  $w$  to a bin  $j$  with load  $L$ 
  - ▶ After this step, bin  $j$  has load  $L + w$
  - ▶ Since  $L$  is the minimum load before the step, the total weight of the items is at least  $mL$  and hence  $L \leq L^*$
  - ▶ Since item  $i$  needs to be placed in some bin, we have  $w \leq L^*$
  - ▶ Hence  $L + w \leq 2L^*$
  - ▶ Thus algorithm  $\mathcal{A}$  achieves an approximation ratio of 2

## Algorithm $\mathcal{A}$ : A Lower Bound of $2 - \varepsilon$

- ▶ Consider an instance with  $n = k(k + 1) + 1$  items and  $m = k + 1$  bins where  $w_i = 1$  for  $1 \leq i \leq k(k + 1)$  and  $w_n = k + 1$
- ▶ After algorithm  $\mathcal{A}$  has processed the first  $k(k + 1)$  items, each bin has a load of  $k$ 
  - ▶ Thus, after the last step, the max bin load is  $2k + 1$
- ▶ It is possible to achieve a max bin load of  $k + 1$ 
  - ▶ Use one bin to hold the item of weight  $k + 1$
  - ▶ Use each of the remaining  $k$  bins to hold  $k + 1$  unit-weight items
- ▶ For any constant  $\varepsilon > 0$ , we can ensure that the approximation ratio exceeds  $2 - \varepsilon$  by choosing  $k$  sufficiently large

# Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{3}{2}$

- ▶ Fix an instance of the load balancing problem, and let  $L^*$  denote the minimax bin load
- ▶ Suppose that at some step of algorithm  $\mathcal{B}$ , we add an item  $i$  with weight  $w$  to a bin  $j$  with load  $L$ 
  - ▶ After this step, bin  $j$  has load  $L + w$
  - ▶ If  $L = 0$ , then  $L + w = w \leq L^*$
  - ▶ Now assume  $L > 0$ 
    - ▶ There are at least  $m + 1$  jobs with weight at least  $w$
    - ▶ Hence  $2w \leq L^*$ , so  $w \leq L^*/2$
    - ▶ Since  $L \leq L^*$ , we conclude that  $L + w \leq (3/2)L^*$

# Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$

- ▶ Fix an instance of the load balancing problem, and let  $L^*$  denote the minimax bin load
- ▶ Assume without loss of generality that  $w_1 \geq \dots \geq w_n$
- ▶ We say that item  $i$  is heavy if  $w_i > L^*/3$ , and is light otherwise
- ▶ Let  $h$  denote the number of heavy items
- ▶ Claim 1: No optimal solution puts more than two heavy items into the same bin
  - ▶ The load of a bin with three heavy items exceeds  $L^*$
- ▶ Claim 2:  $h \leq 2m$ 
  - ▶ Immediate from Claim 1

## Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ Partition the execution of algorithm  $\mathcal{B}$  into a “heavy phase”, in which we place the heavy items  $w_1, \dots, w_h$ , and a “light phase”, in which we place the remaining items
- ▶ Suppose that at some step of algorithm  $\mathcal{B}$  during the light phase, we add an item  $i$  with weight  $w$  to a bin  $j$  with load  $L$ 
  - ▶ Since  $L \leq L^*$  and  $w \leq L^*/3$ , we have  $L + w \leq (4/3)L^*$
- ▶ It remains to consider the heavy phase
  - ▶ Let  $L^{**}$  denote the minimax load for the subinstance with  $m$  bins and only the  $h$  heavy items
  - ▶ Thus  $L^{**} \leq L^*$



## Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ In what follows, we complete the proof by establishing the following result
- ▶ Lemma 1: At the end of the heavy phase, the max bin load is  $L^{**}$
- ▶ By Claim 2,  $h \leq 2m$
- ▶ If  $h \leq m$ , it is easy to see that the claim of Lemma 1 holds
- ▶ For the remainder of the proof, assume that  $h = m + s$  where  $1 \leq s \leq m$
- ▶ We partition the  $h$  heavy items into three sets
  - ▶ The set  $A$  of items 1 through  $m - s$
  - ▶ The set  $B$  of items  $m - s + 1$  through  $m$
  - ▶ The set  $C$  of items  $m + 1$  through  $m + s$

## Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ Let  $\mathcal{X}^*$  denote the set of all ways to assign the heavy items to the bins such that the max bin load is  $L^{**}$ 
  - ▶ These are the optimal solutions to the  $m$ -bin subinstance corresponding to the heavy items  $w_1, \dots, w_h$
  - ▶ We know that no such assignment places more than two items in the same bin
- ▶ Let  $\mathcal{X}$  denote the set of all ways to assign the heavy items to the bins such that there are at most two items in each bin
- ▶ Our plan is to use a series of exchange arguments to prove that the assignment produced by the heavy phase of algorithm  $\mathcal{B}$  belongs to  $\mathcal{X}^*$

## Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ We say that an assignment in  $\mathcal{X}$  is  $A$ -nice if for any item  $i$  in  $A$ , no other item is assigned to the same bin as item  $i$
- ▶ Claim 3: There is an  $A$ -nice assignment in  $\mathcal{X}^*$ 
  - ▶ Let  $\sigma$  be an assignment in  $\mathcal{X}^*$
  - ▶ If  $\sigma$  is not  $A$ -nice, then some item  $i$  in  $A$  shares a bin with another item  $i'$ , while some item  $i''$  in  $B \cup C$  has its own bin
  - ▶ The max load does not increase if we modify  $\sigma$  by moving item  $i'$  to the bin with item  $i''$
  - ▶ This modification reduces by 1 the number of items in  $A$  that share a bin with another item
  - ▶ We can repeat this argument until we reach an  $A$ -nice assignment in  $\mathcal{X}^*$

## Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ We say that an assignment in  $\mathcal{X}$  is  $B$ -nice if no two items in  $B$  are assigned to the same bin
- ▶ Claim 4: There is an assignment in  $\mathcal{X}^*$  that is  $A$ -nice and  $B$ -nice
  - ▶ By Claim 3, there is an  $A$ -nice assignment  $\sigma$  in  $\mathcal{X}^*$
  - ▶ If  $\sigma$  is not  $B$ -nice, then there are items  $i_0$  and  $i_1$  in  $B$  that share a bin, and items  $i_2$  and  $i_3$  in  $C$  that share a bin
  - ▶ The max load does not increase if we modify  $\sigma$  by exchanging items  $i_1$  and  $i_3$ ; moreover, the modified assignment is  $A$ -nice
  - ▶ We can repeat this argument until we reach an assignment in  $\mathcal{X}^*$  that is  $A$ -nice and  $B$ -nice

## Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ Observe that if an assignment in  $\mathcal{X}$  is  $A$ -nice and  $B$ -nice, then each of the  $s$  items in  $B$  shares a bin with one of the  $s$  items in  $C$
- ▶ We say that an assignment in  $\mathcal{X}$  is nice if it is  $A$ -nice,  $B$ -nice, and for every pair of items  $i$  and  $i'$  in  $B$  such that  $w_i > w_{i'}$ , we have  $w_j \leq w_{j'}$  where item  $j$  (resp.,  $j'$ ) in  $C$  shares a bin with item  $i$  (resp.,  $i'$ )
- ▶ Claim 4: Any two nice assignments in  $\mathcal{X}$  have the same distribution of bin loads
  - ▶ Each item in  $A$  has its own bin
  - ▶ Each of the remaining bins contains one item from  $B$
  - ▶ The  $i$ th heaviest item in  $B$  is paired with the  $i$ th lightest item in  $C$  (under some tie breaking) for  $1 \leq i \leq s$

# Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ Claim 5: There is a nice assignment in  $\mathcal{X}^*$ 
  - ▶ By Claim 3, there is an assignment  $\sigma$  in  $\mathcal{X}^*$  that is  $A$ -nice and  $B$ -nice
  - ▶ If  $\sigma$  is not nice, there are items  $i$  and  $i'$  in  $B$  such that  $w_i > w_{i'}$  and  $w_j > w_{j'}$  where item  $j$  (resp.,  $j'$ ) in  $C$  shares a bin with item  $i$  (resp.,  $i'$ )
  - ▶ The max load does not increase if we modify  $\sigma$  by exchanging items  $j$  and  $j'$ ; moreover, the modified assignment is  $A$ -nice and  $B$ -nice
  - ▶ We can repeat this argument until we reach a nice assignment in  $\mathcal{X}^*$

## Algorithm $\mathcal{B}$ : An Upper Bound of $\frac{4}{3}$ (cont'd)

- ▶ It is easy to see that algorithm  $\mathcal{B}$  produces a nice assignment in  $\mathcal{X}$  at the end of the heavy phase
- ▶ By Claims 4 and 5, we conclude that  $\mathcal{B}$  produces an assignment in  $\mathcal{X}^*$  at the end of the heavy phase
- ▶ This completes the proof of Lemma 1

# Algorithm $\mathcal{B}$ : A Lower Bound of $\frac{4}{3} - \varepsilon$

- ▶ Exercise: As in the  $2 - \varepsilon$  lower bound argument for algorithm  $\mathcal{A}$ , identify a suitable family of “bad” instances for algorithm  $\mathcal{B}$