

# An Activity Selection Problem

- ▶ We are given  $n$  activities indexed from 1 to  $n$
- ▶ Activity  $i$  has start time  $s_i$  and finish time  $f_i$
- ▶ We cannot participate in two activities that overlap in time
- ▶ We wish to determine a maximum-cardinality set of non-overlapping activities

# Key Observation

- ▶ Let  $i$  be an activity with minimum finish time
- ▶ Claim: Some optimal solution includes  $i$
- ▶ To prove this claim, we can use an “exchange argument”
  - ▶ Suppose  $S$  is an optimal solution that does not include  $i$
  - ▶ Let  $j$  be the first activity in  $S$
  - ▶ Then  $(S - j) + i$  is an optimal solution that includes  $i$

# A Greedy Algorithm

- ▶ Re-index the activities in nondecreasing order of finish time
- ▶ Initialize  $I$  to  $\{1, \dots, n\}$  and  $S$  to  $\emptyset$
- ▶ While  $I \neq \emptyset$ 
  - ▶ Let  $i$  be a minimum-index activity in  $I$
  - ▶ Add  $i$  to  $S$
  - ▶ Eliminate from  $I$  all indices  $j$  such that activities  $i$  and  $j$  overlap

# A Fast Implementation

- ▶ Re-index the activities in nondecreasing order of finish time
- ▶ Initialize  $S$  to  $\{1\}$  and  $k$  to 1
- ▶ For  $i$  running from 2 to  $n$ 
  - ▶ If the start time of activity  $i$  is at least the finish time of activity  $k$ , then add  $i$  to  $S$  and set  $k$  to  $i$

# The Fractional Knapsack Problem

- ▶ Recall the knapsack problem
  - ▶ We are given a positive integer knapsack capacity  $W$  and  $n$  items indexed from 1 to  $n$
  - ▶ Item  $i$  has positive integer value  $v_i$  and weight  $w_i$
  - ▶ We wish to identify a maximum-value set of items with weight at most  $W$
- ▶ In the fractional knapsack problem, we are allowed to take a fractional amount of any item

# Key Observation

- ▶ Let  $i$  be an item with maximum “value density”  $v_i/w_i$
- ▶ Claim: Some optimal solution includes a  $z = \min(1, W/w_i)$  fraction of item  $i$
- ▶ To prove this claim, we can use an exchange argument
  - ▶ Let  $S$  be an optimal solution that includes a fraction  $z' < z$  of item  $i$
  - ▶ Observe that the weight of  $S$  is at least  $z$
  - ▶ Modify  $S$  by removing (fractional) items not equal to  $i$  with total weight  $z - z'$ , and replacing them with  $z - z'$  units of item  $i$
  - ▶ Observe that  $S$  remains optimal

# A Greedy Algorithm

- ▶ Re-index the items in nonincreasing order of value density
- ▶ Take as much as possible of item 1, then as much as possible of item 2, et cetera, until the knapsack is full or there are no items left
- ▶ This algorithm uses  $O(n \log n)$  operations due to the sorting (re-indexing) step
- ▶ Can we do better?

# A Faster Implementation

- ▶ Recall the BFPRT linear-time selection algorithm
- ▶ It is easy to generalize the BFPRT algorithm to solve the following weighted selection problem in linear time
  - ▶ Each of the  $n$  keys in the input has a positive weight  $w_i$
  - ▶ If some keys are equal, choose a way to break ties (e.g., by index) to obtain a total ordering of the keys
  - ▶ Let  $X$  denote  $\sum_{1 \leq i \leq n} w_i$
  - ▶ We are given a desired threshold  $x$ ,  $0 \leq x \leq X$
  - ▶ We wish to identify the maximum key  $k$  such that the total weight of the keys preceding key  $k$  is less than  $x$
- ▶ We can use a linear-time weighted selection algorithm to solve fractional knapsack in linear time



# Scheduling to Minimize Maximum Lateness

- ▶ We are given  $n$  tasks indexed from 1 to  $n$
- ▶ Task  $i$  has a positive integer deadline  $d_i$  and a positive integer execution requirement  $e_i$
- ▶ We wish to (nonpreemptively) schedule all  $n$  tasks on a single resource beginning at time 0 in such a way that the maximum “lateness” of any task is minimized
  - ▶ A task with deadline  $d$  and termination time  $t$  is defined to have lateness  $\max(0, t - d)$
- ▶ We can restrict attention to gap-free schedules, so we are optimizing over  $n!$  schedules

# Key Lemma

- ▶ Suppose  $S$  is a schedule in which task  $j$  is executed immediately after task  $i$  and  $d_j \leq d_i$ 
  - ▶ Let  $\ell_i$  (resp.,  $\ell_j$ ) denote the lateness of task  $i$  (resp.,  $j$ ) in  $S$
- ▶ Let  $S'$  be the schedule that is the same as  $S$  except that the order of execution of tasks  $i$  and  $j$  is interchanged
  - ▶ Let  $\ell'_i$  (resp.,  $\ell'_j$ ) denote the lateness of task  $i$  (resp.,  $j$ ) in  $S'$
- ▶ Lemma:  $\ell_j \geq \max(\ell'_i, \ell'_j)$
- ▶ This lemma implies that the “earliest deadline” rule yields an optimal schedule
  - ▶ Ties can be broken arbitrarily

# Proof of the Key Lemma

- ▶ Lemma:  $\ell_j \geq \max(\ell'_i, \ell'_j)$ 
  - ▶ Assume tasks  $i$  and  $j$  are executed in the time interval  $[s, s + e_i + e_j]$  in  $S$  and  $S'$
  - ▶ We have  $\ell_j = \max(0, A)$  where  $A = s + e_i + e_j - d_j$ ,  $\ell'_i = \max(0, B)$  where  $B = s + e_i + e_j - d_i$ , and  $\ell'_j = \max(0, C)$  where  $C = s + e_j - d_j$
  - ▶ Observe that  $A > C$
  - ▶ Since  $d_j \leq d_i$ , we have  $A \geq B$
  - ▶ The lemma follows since

$$\ell_j = \max(0, A) \geq \max(0, B, C) = \max(\ell'_i, \ell'_j)$$

# Single-Source Shortest Paths, Revisited

- ▶ Recall that we can solve the SSSP problem in  $O(|E| \cdot |V|)$  using the Bellman-Ford algorithm, which can handle negative edge weights
- ▶ If the edge weights are nonnegative, we can solve the SSSP problem much more rapidly using Dijkstra's algorithm
- ▶ For any vertex  $v$ , let  $d(v)$  denote the shortest path distance from the source  $s$  to  $v$ 
  - ▶ Thus  $d(s) = 0$
  - ▶ We will maintain a “label”  $\ell_v$  for each vertex  $v$  in  $V$
  - ▶ We initialize  $\ell_s$  to 0 and  $\ell_v$  to  $\infty$  for  $v \neq s$

# Dijkstra's SSSP Algorithm

- ▶ We maintain a subset  $U$  of  $V$  that is initialized to  $\{s\}$
- ▶ In each of  $|V| - 1$  iterations, we add a vertex to  $U$
- ▶ We maintain the following key invariants
  - ▶ For each vertex  $u$  in  $U$ , we have  $\ell_u = d(u)$
  - ▶ For each vertex  $v$  in  $V \setminus U$ , we have

$$\ell_v = \min_{u \in U: (u,v) \in E} d(u) + w(u, v)$$

(or  $\infty$  if the minimization is over an empty set); this is an upper bound on  $d(v)$

- ▶ How do we choose which vertex to add to  $U$  in each iteration, and how do we maintain the key invariants?

# A General Iteration

- ▶ Let  $u$  be a minimum-label vertex in  $V \setminus U$ 
  - ▶ Observe that any path from  $s$  to a vertex in  $V \setminus U$  has cost at least  $\ell_u$ ; hence  $\ell_u \leq d(u)$
  - ▶ Since  $d(u) \leq \ell_u$  by the second key invariant, we conclude that  $\ell_u = d(u)$
- ▶ We add  $u$  to  $U$
- ▶ The first key invariant is maintained since  $\ell_u = d(u)$
- ▶ To re-establish the second key invariant, we update  $\ell_v$  to

$$\min(\ell_v, d(u) + w(u, v))$$

for each vertex  $v$  in  $V \setminus U$  such that  $(u, v) \in E$

# Efficient Implementation of Dijkstra's Algorithm

- ▶ We can use a heap to maintain the labels of the vertices in  $V \setminus U$ 
  - ▶ We use  $O(|V|)$  INSERT and DELETE-MIN operations
  - ▶ We use  $O(|E|)$  DECREASE-KEY operations
- ▶ Using an elementary heap data structure, the algorithm runs in  $O(|E| \log |V|)$  time
  - ▶ Using an array to maintain the labels, we obtain an  $O(|V|^2)$  bound, which is an improvement for sufficiently dense graphs
- ▶ Using a more sophisticated data structure such as a Fibonacci heap (to be discussed in a later lecture), we obtain a bound of  $O(|E| + |V| \log |V|)$

# The Minimum Spanning Tree Problem

- ▶ We are given a connected, undirected graph  $G = (V, E)$  where each edge  $e$  in  $E$  has an associated weight  $w(e)$  (which may be negative)
- ▶ A (graph-theoretic) tree is a graph that is acyclic and connected
- ▶ A spanning tree  $T$  of  $G$  is a subgraph  $G' = (V, E')$  of  $G$  that is a tree
  - ▶ It is convenient to identify  $T$  with its edge set
  - ▶ It is easy to prove that all spanning trees of  $G$  have cardinality  $|V| - 1$
- ▶ The weight of a spanning tree  $T$  is defined as  $\sum_{e \in T} w(e)$
- ▶ A minimum spanning tree (MST) of  $G$  is a spanning tree of  $G$  of minimum weight



# Key Observation

- ▶ Let  $e$  be a minimum-weight edge in  $E$
- ▶ Claim 1: Some MST of  $G$  includes  $e$
- ▶ To prove this claim, we can use an exchange argument
  - ▶ Let  $T$  be an MST of  $G$  that does not include  $e$
  - ▶ If we add  $e$  to  $T$  we get a unique cycle  $C$
  - ▶ For any edge  $e'$  on  $C$ ,  $T + e - e'$  is a spanning tree of  $G$  with weight at most that of  $T$ , and hence is an MST of  $G$
- ▶ We can obtain an MST of  $G$  by recursively computing an MST  $T'$  of the graph  $G'$  obtained by “contracting” edge  $e$ , and returning  $T' + e$ 
  - ▶ Alternatively, we can get an iterative implementation by repeatedly selecting a minimum-weight edge that does not form a cycle with any subset of the previously selected edges

# Kruskal's MST Algorithm

- ▶ Index the edges  $e_1, \dots, e_{|E|}$  in nondecreasing order of weight
- ▶ Initialize  $T$  to  $\emptyset$
- ▶ For  $i$  running from 1 to  $|E|$ 
  - ▶ If  $T + e_i$  is acyclic, then add  $e_i$  to  $T$
- ▶ Return  $T$

# Correctness of Kruskal's Algorithm

- ▶ We first claim that the output  $T$  is a spanning tree of  $G$
- ▶ Clearly,  $T$  is acyclic
- ▶ Suppose  $(V, T)$  has connected components  $G_1, \dots, G_k$  where  $k > 1$
- ▶ Since  $G$  is connected, there is an edge  $e$  in  $E$  such that the two endpoints of  $e$  belong to distinct components  $G_i$  and  $G_j$
- ▶ But then  $T + e$  is acyclic, a contradiction

# Correctness of Kruskal's Algorithm (cont'd)

- ▶ It remains to prove that  $T$  is an MST of  $G$
- ▶ Let  $T_0$  be an MST of  $G$
- ▶ If  $T = T_0$ , we are done, so assume  $T \neq T_0$ 
  - ▶ Let  $i$  be the least integer such that  $e_i \in T \oplus T_0$
  - ▶ Observe that  $e_i \in T \setminus T_0$
  - ▶ Let  $C$  denote the unique cycle in  $T_0 + e_i$
  - ▶ Let  $e_j$  denote an edge in  $C \cap (T_0 \setminus T)$ ; thus  $i < j$  and  $w(e_i) \leq w(e_j)$
  - ▶ Since  $T_0$  is an MST and  $T_1 = T_0 + e_i - e_j$  is a spanning tree, we have  $w(e_j) \leq w(e_i)$
  - ▶ Thus  $w(e_i) = w(e_j)$  and  $T_1$  is an MST with  $|T \cap T_1| = |T \cap T_0| + 1$

# Correctness of Kruskal's Algorithm (cont'd)

- ▶ If  $T_1 \neq T$ , we can repeat the previous argument with  $T_1$  playing the role of  $T_0$  to obtain an MST  $T_2$  such that  $|T \cap T_2| = |T \cap T_0| + 2$
- ▶ Let  $k$  denote  $|V| - 1 - |T \cap T_0|$
- ▶ After  $k$  iterations, we obtain a sequence of  $k + 1$  MSTs  $T_0, \dots, T_k$  such that  $T = T_k$ 
  - ▶ Thus  $T$  is an MST
- ▶ All of the  $T_i$ 's have the same distribution of edge weights
  - ▶ That is, for any given weight  $z$ , all of the  $T_i$ 's contain the same number of edges of weight  $z$
  - ▶ Since  $T_0$  is an arbitrary MST of  $G$ , we conclude that all MSTs of  $G$  have the same distribution of edge weights

# Some Other Properties of MSTs

- ▶ If all of the edge weights are distinct, there is a unique MST
  - ▶ Follows from the preceding claim that all MSTs have the same distribution of edge weights
- ▶ Kruskal's algorithm can generate any MST
  - ▶ When indexing the edges, Kruskal's algorithm can break ties arbitrarily
  - ▶ To produce MST  $T$ , favor edges in  $T$  over edges not in  $T$
- ▶ The set of MSTs does not change if we replace each weight  $x$  with  $f(x)$  for some increasing function  $f$ 
  - ▶ This transformation preserves the relative order of the weights, which is all that Kruskal's algorithm looks at

# Efficient Implementation of Kruskal's Algorithm

- ▶ Adding an edge  $e$  to  $T$  creates a cycle if and only if the two endpoints of  $e$  belong to the same connected component of  $(V, T)$
- ▶ If  $T + e$  is acyclic, so that we add  $e$  to  $T$ , then the two connected components bridged by  $e$  are merged into one
- ▶ We can use a “union-find” data structure to manage the vertex sets associated with the connected components of  $(V, T)$

# Union-Find

- ▶ A union-find data structure maintains a collection of disjoint sets, subject to the following operations
- ▶  $\text{MAKE-SET}(x)$  forms a new singleton set  $\{x\}$ 
  - ▶ To maintain disjointness, we require that  $x$  does not belong to any of the existing sets in the collection
- ▶  $\text{UNION}(x, y)$  merges the set containing  $x$  with the set containing  $y$ , where  $x$  and  $y$  belong to distinct sets
- ▶  $\text{FIND-SET}(x)$  returns the “name” of the set containing  $x$  ( $x$  is required to belong to some set)
  - ▶ We require that  $\text{FIND-SET}(x) = \text{FIND-SET}(y)$  if and only if  $x$  and  $y$  belong to the same set
- ▶ We will study a fast union-find data structure in a later lecture



# A Union-Find Implementation of Kruskal's Algorithm

- ▶ At the outset, we perform a  $\text{MAKE-SET}(v)$  operation for each  $v$  in  $V$ 
  - ▶ Thus our initial sets are the vertex sets of the  $|V|$  connected components of  $(V, \emptyset)$
- ▶ For an edge  $e = (u, v)$ , we check whether  $T + e$  contains a cycle by asking whether  $\text{FIND-SET}(u) = \text{FIND-SET}(v)$ 
  - ▶ We perform  $2|E|$   $\text{FIND-SET}$  operations
- ▶ When we add an edge  $e = (u, v)$  to  $T$ , we perform a  $\text{UNION}(u, v)$  operation
  - ▶ We perform  $|V| - 1$   $\text{UNION}$  operations