

Matching in Regular Bipartite Graphs

- ▶ For any positive integers n and d , let $\mathcal{G}_{n,d}$ denote the set of all d -regular bipartite graphs (U, V, E) where $|U| = |V| = n$ and $E \subseteq U \times V$
- ▶ Using Hall's theorem, it is easy to argue that any bipartite graph in $\mathcal{G}_{n,d}$ admits a perfect matching
 - ▶ Indeed, the edge set can be partitioned into d perfect matchings
- ▶ We will present a fast randomized algorithm for computing a perfect matching in a given bipartite graph in $\mathcal{G}_{n,d}$

High-Level Description of the Algorithm

- ▶ Recall the elementary flow-based approach to computing a maximum cardinality matching of an arbitrary bipartite graph
 - ▶ Repeatedly augment a 0-1 flow by one unit
 - ▶ A 0-1 flow of value ℓ corresponds to a matching of cardinality ℓ
 - ▶ For a d -regular bipartite graph, the running time is $O(dn)$ per iteration for an overall running time of $O(dn^2)$
- ▶ Our plan is to use a random walk to find an augmenting path
 - ▶ If k vertices on each side remain unmatched, we will show that the expected time for the random walk to find an augmenting path is $O(n/k)$
 - ▶ This results in an overall expected running time of $O(n \log n)$

Analysis of a Single Augmentation: Basic Definitions

- ▶ Assume that our bipartite graph $G = (U, V, E)$ belongs to $\mathcal{G}_{n,d}$
- ▶ Let U_0 (resp., U_1) denote the unmatched (resp., matched) vertices in U
- ▶ Let V_0 (resp., V_1) denote the unmatched (resp., matched) vertices in V
- ▶ Let μ denote the current matching
 - ▶ For any vertex u in U_1 , we write $\mu(u)$ to denote the match of u in V_1
 - ▶ For any vertex v in V_1 , we write $\mu(v)$ to denote the match of v in U_1

The Random Walk

- ▶ Instead of introducing a source and sink, we will start at a uniformly random node in U_0 , and we will terminate once we reach any node in V_0
 - ▶ From a node u in U_0 , we go to a uniformly random node in $\Gamma(u)$
 - ▶ From a node u in U_1 , we go to a uniformly random node in $\Gamma(u) - \mu(u)$
 - ▶ From a node v in V_1 , we go to $\mu(v)$
- ▶ We can easily prune the random walk to obtain a simple alternating path for augmenting the matching μ
- ▶ For any node x , let $b(x)$ denote the expected number of “back” edges (i.e., edges from V to U) traversed before the walk terminates, assuming that we are currently at node x

- ▶ Lemma 1: For any v in V , we have

$$b(v) = \begin{cases} 0 & \text{if } v \in V_0 \\ 1 + b(\mu(v)) & \text{if } v \in V_1 \end{cases}$$

- ▶ If v belongs to V_0 then the walk terminates
- ▶ If v belongs to V_1 then we traverse the back edge from v to $\mu(v)$, and the walk continues from $\mu(v)$

- ▶ Lemma 2: For any u in U_0 , we have

$$d \cdot b(u) = \sum_{v \in \Gamma(u)} b(v)$$

- ▶ From a vertex u in U_0 , we select a uniformly random vertex v in $\Gamma(u)$, traverse the “forward” edge from u to v , and then continue the walk from v
- ▶ Hence $b(u)$ is equal to $\frac{1}{d} \sum_{v \in \Gamma(u)} b(v)$

- ▶ Lemma 3: For any u in U_1 , we have

$$d \cdot b(u) = -1 + \sum_{v \in \Gamma(u)} b(v)$$

- ▶ From a vertex u in U_1 , we select a uniformly random vertex v in $\Gamma(u) - \mu(u)$, traverse the “forward” edge from u to v , and then continue the walk from $\mu(v)$
- ▶ Hence $b(u)$ is equal to $[-b(\mu(u)) + \sum_{v \in \Gamma(u)} b(v)]/(d - 1)$
- ▶ Equivalently, $(d - 1)b(u) = -b(\mu(u)) + \sum_{v \in \Gamma(u)} b(v)$
- ▶ The claim follows since $b(\mu(u)) = 1 + b(u)$

- ▶ Lemma 4: $d \sum_{u \in U} b(u) = (d-1)|\mu| + d \sum_{u \in U_1} b(u)$
 - ▶ By Lemmas 2 and 3, the LHS is equal to $-|\mu| + \sum_{u \in U} \sum_{v \in \Gamma(u)} b(v)$
 - ▶ Since the graph is d -regular, this is equal to $-|\mu| + d \sum_{v \in V} b(v)$
 - ▶ By Lemma 1, this is equal to $-|\mu| + d(|\mu| + \sum_{u \in U_1} b(u))$
 - ▶ The claim follows

Analysis (cont'd)

- ▶ Let T_k denote the expected length of the random walk when $|U_0| = |V_0| = k$ and hence $|\mu| = n - k$
- ▶ Thus $T_k = 1 + \frac{2}{k} \sum_{u \in U_0} b(u)$
- ▶ Lemma 4 implies that $d \sum_{u \in U_0} b(u) = (d - 1)(n - k)$
- ▶ Hence $T_k = 1 + \frac{2(d-1)(n-k)}{dk} < 1 + 2n/k$
- ▶ Thus the overall expected running time is

$$O\left(\sum_{1 \leq k \leq n} T_k\right) = O(nH_n) = O(n \log n)$$

How Robust is this Result?

- ▶ Intuitively, one might expect the $O(n \log n)$ expected time bound to hold for graphs that are “nearly” regular (and contain a perfect matching, say)
- ▶ Somewhat surprisingly, even a small deviation from regularity can dramatically worsen the expected running time
- ▶ We will demonstrate a bad example where we add a single additional edge to a 3-regular bipartite graph
- ▶ We begin by discussing a random walk on a cycle

A Random Walk on a Cycle

- ▶ Consider a cycle with n nodes indexed from 0 to $n - 1$ in clockwise order
- ▶ Consider a random walk starting at node i
 - ▶ At each step of the walk, we use an independent fair coin flip to determine whether to move one step clockwise or one step counterclockwise
 - ▶ The walk ends when we reach node 0
 - ▶ It is not hard to prove that the expected number of steps is $i(n - i)$
 - ▶ Thus if we start at $\Theta(n)$ distance from node 0, the expected running time is $\Theta(n^2)$

A Random Walk on a Cycle (cont'd)

- ▶ Suppose that we start the random walk at a node $i \neq 0$
- ▶ Let E denote the event that the random walk terminates by moving from node 1 to node 0
- ▶ Thus $\neg E$ is the event that the random walk terminates by moving from node $n - 1$ to node 0
- ▶ It is not hard to prove that $\Pr(E) = (n - i)/n$

A Bad Example

- ▶ Let n be an even positive integer, and consider the following bipartite graph $G = (U, V, E)$ in $\mathcal{G}_{n,3}$
 - ▶ Index the nodes in U (resp. V) from 0 to $n - 1$
 - ▶ Connect each node i in U to nodes $(i - 1) \bmod n$, i , and $(i + 1) \bmod n$ in V
- ▶ Construct $G' = (U, V, E')$ from G by setting E' to $E + (u, v)$ where u is the node with index 0 in U and v is node the node with index $n/2$ in V
- ▶ Assume that our current matching μ matches node i in U to node i in V for $1 \leq i \leq n$
- ▶ What is the expected running time of the random walk in G' that starts at the node with index 0 in U and terminates at the node with index 0 in V ?

A Bad Example (cont'd)

- ▶ Starting at the node with index 0 in U , there is a $1/4$ chance that we move to the node with index $n/2$ in V in the first step
- ▶ If we do this, then based on our discussion of the random walk on a cycle, the expected number of steps to reach the node with index 0 in V is $\Omega(n^2)$

A “High Probability” Result?

- ▶ Let us define “with high probability” to mean “with arbitrary inverse polynomial failure probability”
- ▶ Does the random walk algorithm have $O(n \log n)$ running time with high probability?
- ▶ Consider the 3-regular graph G associated with the previous construction
- ▶ Based on our discussion of the random walk on a cycle, there is an $\Omega(1/n)$ chance that the random walk on G will reach the node with index $n/2$ in V
- ▶ Thus the augmentation takes $\Theta(n^2)$ time with probability $\Omega(1/n)$

An “Abort-and-Restart” Variant

- ▶ A natural variant of the random walk algorithm does achieve $O(n \log n)$ time with high probability
 - ▶ Suppose $|U_0| = |V_0| = k$
 - ▶ Recall that the expected length of the random walk is less than $1 + 2n/k$
 - ▶ Markov's inequality implies that the random walk terminates within $2(1 + 2n/k)$ steps with probability at least $1/2$
 - ▶ If the random walk fails to terminate within $2(1 + 2n/k)$ steps, we restart it from the beginning

Analysis of the Abort-and-Restart Variant

- ▶ We claim that the overall running time of the abort-and-restart variant of the random walk algorithm is $O(n \log n)$ with high probability
- ▶ To establish this, we instead analyze a related game
 - ▶ It is straightforward to verify that the bound we establish for the game implies the claim

The Game

- ▶ A video game has $\log_2 n$ “levels” indexed from 1 to $\log_2 n$, where n is a power of 2
- ▶ You start play at level 1 with “energy” $cn \log_2 n$ and repeatedly fight against one monster at a time
- ▶ Each time you fight with a monster on any level, you have a fifty percent chance of defeating the monster
- ▶ To advance from level ℓ to level $\ell + 1$, you need to defeat $n2^{-\ell}$ level- ℓ monsters
- ▶ Whenever you fight with a level- ℓ monster, your energy is reduced by 2^ℓ , whether or not you defeat the monster
- ▶ You win if you complete all of the levels without running out of energy

Analysis of the Game

- ▶ The expected number of level- ℓ monsters faced is $2n2^{-\ell}$
- ▶ By Markov's inequality, the probability that you need to fight more than $4n2^{-\ell}$ level- ℓ monsters is at most $\frac{1}{2}$
- ▶ The cost of facing $4n2^{-\ell}$ level- ℓ monsters is $4n$
- ▶ Let the random variable Z denote the number of independent flips of a fair coin required to get $\log_2 n$ heads
- ▶ Observe that the probability you lose the game is upper bounded by the probability that $4nZ \geq cn \log_2 n$, i.e., that $Z \geq (c/4) \log_2 n$

Analysis of the Game (cont'd)

- ▶ The probability that Z exceeds $(c/4) \log_2 n$ is equal to the probability that we get fewer than $\log_2 n$ heads in $(c/4) \log_2 n$ independent flips of a fair coin
- ▶ Recall that for a random variable X drawn from $B(k, 1/2)$ and any δ in $[0, 1]$, we have

$$\Pr(X \leq (1 - \delta)k/2) \leq \exp(-\delta^2 k/2)$$

- ▶ Using the above bound, the probability Z exceeds $(c/4) \log_2 n$ is upper bounded by an inverse polynomial in n where the exponent is quadratic in c
- ▶ Thus the probability you lose the game is upper bounded by an arbitrary inverse polynomial for a sufficiently large choice of the constant c